Copyleaks

Analysis Report

Plagiarism Detection and AI Detection Report

New Scans 9:34 AM

Scan details

Scan Time	Total Pages	Total Words
April 8th, 2026 at 09:34 UTC	3	575

Plagiarism Detection

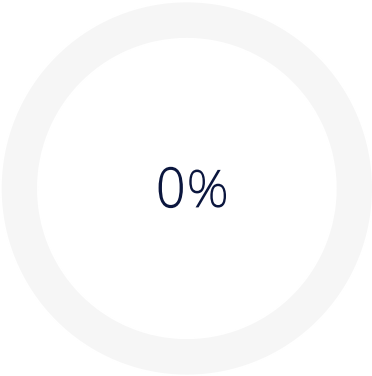

Plagiarism types	Text coverage	Words
Identical	0%	0
Minor Changes	0%	0
Paraphrased	0%	0
Excluded		
Omitted Words		0

AI Detection

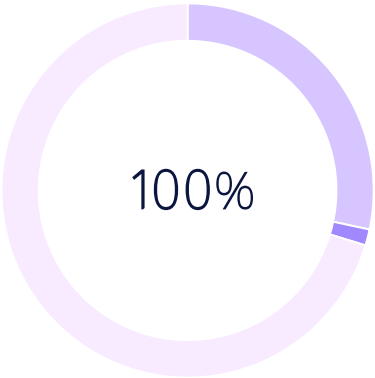

	Text coverage	Words
AI Text	100%	575
Low Frequency		163
Medium Frequency		0
High Frequency		8
Human Text	0%	0
Excluded		
Omitted Words		0

Plagiarism

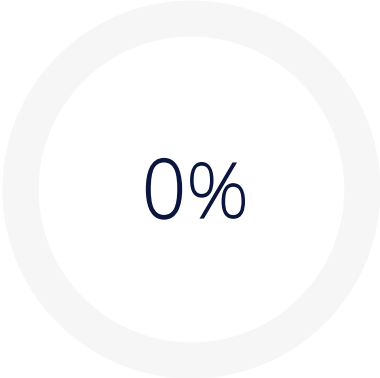

Results (0)

*Results may not appear because the feature has been disabled.

Private Cloud Hub	Shared Data Hub	Filtered / Excluded
0	0	0
Internet Sources	**AI Source Match**	**Current Batch**
0	0	0

Plagiarism types	Text coverage	Words
Identical	0%	0
Minor Changes	0%	0
Paraphrased	0%	0
Excluded		
Omitted Words		0

About Plagiarism Detection

Our AI-powered plagiarism scans offer three layers of text similarity detection: Identical, Minor Changes, and Paraphrased. Based on your scan settings we also provide insight on how much of the text you are not scanning for plagiarism (Omitted words).

Identical

One to one exact word matches. Learn more

Minor Changes

Words that hold nearly the same meaning but have a change to their form (e.g. “large” becomes “largely”). Learn more

Paraphrased

Different words that hold the same meaning that replace the original content (e.g. 'large' becomes 'big') Learn more

Omitted Words

The portion of text that is not being scanned for plagiarism based on the scan settings. (e.g. the 'Ignore quotations' setting is enabled and the document is 20% quotations making the omitted words percentage 20%) Learn more

Copyleaks Shared Data Hub

Our Shared Data Hub is a collection of millions of user-submitted documents that you can utilize as a scan resource and choose whether or not you would like to submit the file you are scanning into the Shared Data Hub. Learn more

Filtered and Excluded Results

The report will generate a complete list of results. There is always the option to exclude specific results that are not relevant. Note, by unchecking certain results, the similarity percentage may change. Learn more

Current Batch Results

These are the results displayed from the collection, or batch, of files uploaded for a scan at the same time. Learn more

AI Content

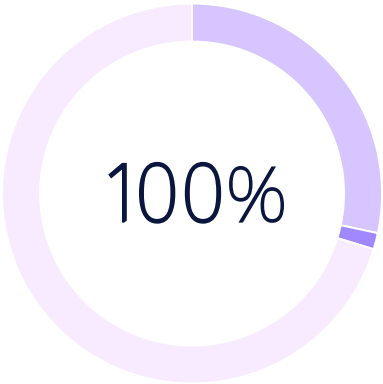

	Text coverage	Words
AI Text	100%	575
Low Frequency		163
Medium Frequency		0
High Frequency		8
Human Text	0%	0
Excluded		
Omitted Words		0

About AI Detection

Our AI Detector is the only enterprise-level solution that can verify if the content was written by a human or generated by AI, including source code and text that has been plagiarized or modified. Learn more

AI Text

A body of text that has been generated or altered by AI technology. Learn more

Human Text

Any text that has been fully written by a human and has not been altered or generated by AI. Learn more

Copyleaks AI Detector Effectiveness

Credible data at scale, coupled with machine learning and widespread adoption, allows us to continually refine and improve our ability to understand complex text patterns, resulting in over 99% accuracy–far higher than any other AI detector–and improving daily. Learn more

Ideal Text Length

The higher the character count, the easier for our technology to determine irregular patterns, which results in a higher confidence rating for AI detection. Learn more

Reasons It Might Be AI When You Think It's Not

The AI Detector can detect a variety of AI-generated text, including tools that use AI technology to paraphrase content, auto-complete sentences, and more. Learn more

User AI Alert History

Historical data of how many times a user has been flagged for potentially having AI text within their content. Learn more

AI Logic

The number of times a phrase was found more frequently in AI vs human text is shown according to low, medium, and high frequency. Learn more

AI Logic

Shows you the "why" behind AI detection with sources you can see and verify.

AI Phrases

Detects phrases that appear with higher frequency in AI-written text than in human writing.

The frequency of a phrase in AI vs. human text.

3 x — > 10,000x

> 10,000x Large Language Model (LLM)

How frequently the phrase was found in our dataset:

AI Text	**1370.14 / 1,000,000 Documents**
Human Text	**0.04 / 1,000,000 Documents**

> 10,000x allowing the LLM to

How frequently the phrase was found in our dataset:

AI Text	**38.06 / 1,000,000 Documents**
Human Text	**0 / 1,000,000 Documents**

5237x the LLM to

How frequently the phrase was found in our dataset:

AI Text	**791.25 / 1,000,000 Documents**
Human Text	**0.15 / 1,000,000 Documents**

3425x the LLM prompt

How frequently the phrase was found in our dataset:

AI Text	**4.38 / 1,000,000 Documents**
Human Text	**0 / 1,000,000 Documents**

2770x outputs are grounded in

How frequently the phrase was found in our dataset:

AI Text	**17.74 / 1,000,000 Documents**
Human Text	**0.01 / 1,000,000 Documents**

1182x The SFT data

How frequently the phrase was found in our dataset:

AI Text	**34.8 / 1,000,000 Documents**
Human Text	**0.03 / 1,000,000 Documents**

1118x LLM to: "Ensure

How frequently the phrase was found in our dataset:

AI Text	**4.3 / 1,000,000 Documents**
Human Text	**0 / 1,000,000 Documents**

1113x reasoning logically leads to the

How frequently the phrase was found in our dataset:

AI Text	**2.85 / 1,000,000 Documents**
Human Text	**0 / 1,000,000 Documents**

929x to Reward Hacking and

How frequently the phrase was found in our dataset:

AI Text	**1.19 / 1,000,000 Documents**
Human Text	**0 / 1,000,000 Documents**

766x the model's capacity for

How frequently the phrase was found in our dataset:

AI Text	**3.92 / 1,000,000 Documents**
Human Text	**0.01 / 1,000,000 Documents**

751x the LLM from

How frequently the phrase was found in our dataset:

AI Text	**22.12 / 1,000,000 Documents**
Human Text	**0.03 / 1,000,000 Documents**

743x The prompt template for

How frequently the phrase was found in our dataset:

AI Text	**3.8 / 1,000,000 Documents**
Human Text	**0.01 / 1,000,000 Documents**

255x with advanced reasoning capabilities

How frequently the phrase was found in our dataset:

AI Text	**3.91 / 1,000,000 Documents**
Human Text	**0.02 / 1,000,000 Documents**

237x into a rigid grid.

How frequently the phrase was found in our dataset:

AI Text	**4.25 / 1,000,000 Documents**
Human Text	**0.02 / 1,000,000 Documents**

233x **Summary: The paper proposes**

How frequently the phrase was found in our dataset:

AI Text	**21.46 / 1,000,000 Documents**
Human Text	**0.09 / 1,000,000 Documents**

227x **and "Teacher Forcing" in**

How frequently the phrase was found in our dataset:

AI Text	**2.32 / 1,000,000 Documents**
Human Text	**0.01 / 1,000,000 Documents**

162x **The authors attribute this to**

How frequently the phrase was found in our dataset:

AI Text	**114.19 / 1,000,000 Documents**
Human Text	**0.7 / 1,000,000 Documents**

100x **component of this framework is**

How frequently the phrase was found in our dataset:

AI Text	**17.99 / 1,000,000 Documents**
Human Text	**0.18 / 1,000,000 Documents**

86x **the empirical validation**

How frequently the phrase was found in our dataset:

AI Text	**106.17 / 1,000,000 Documents**
Human Text	**1.24 / 1,000,000 Documents**

47x **ensures that intermediate**

How frequently the phrase was found in our dataset:

AI Text	**2.76 / 1,000,000 Documents**
Human Text	**0.06 / 1,000,000 Documents**

44x **MDP Formulation and**

How frequently the phrase was found in our dataset:

AI Text	**3.56 / 1,000,000 Documents**
Human Text	**0.08 / 1,000,000 Documents**

36x **explicitly instructs the**

How frequently the phrase was found in our dataset:

AI Text	**7.72 / 1,000,000 Documents**
Human Text	**0.22 / 1,000,000 Documents**

28x **This directly explains the**

How frequently the phrase was found in our dataset:

AI Text	**1.48 / 1,000,000 Documents**
Human Text	**0.05 / 1,000,000 Documents**

20x **While the integration of**

How frequently the phrase was found in our dataset:

AI Text	**29.13 / 1,000,000 Documents**
Human Text	**1.45 / 1,000,000 Documents**

19x **Behavioral Cloning, and**

How frequently the phrase was found in our dataset:

AI Text	**3.36 / 1,000,000 Documents**
Human Text	**0.17 / 1,000,000 Documents**

18x **the World Model is**

How frequently the phrase was found in our dataset:

AI Text	**9.09 / 1,000,000 Documents**
Human Text	**0.49 / 1,000,000 Documents**

17x **likely to hallucinate**

How frequently the phrase was found in our dataset:

AI Text	**1.99 / 1,000,000 Documents**
Human Text	**0.12 / 1,000,000 Documents**

16x **fundamentally flawed due to**

How frequently the phrase was found in our dataset:

AI Text	**3.32 / 1,000,000 Documents**
Human Text	**0.21 / 1,000,000 Documents**

13x **the intermediate reward**

How frequently the phrase was found in our dataset:

AI Text	**1.83 / 1,000,000 Documents**
Human Text	**0.14 / 1,000,000 Documents**

12x **Pipeline Design: The**

How frequently the phrase was found in our dataset:

AI Text	**2.29 / 1,000,000 Documents**
Human Text	**0.19 / 1,000,000 Documents**

11x **World Model" that**

How frequently the phrase was found in our dataset:

AI Text	**9.39 / 1,000,000 Documents**
Human Text	**0.85 / 1,000,000 Documents**

9x **matching, but this**

How frequently the phrase was found in our dataset:

AI Text	**3.36 / 1,000,000 Documents**
Human Text	**0.36 / 1,000,000 Documents**

8x **is outperformed by the**

How frequently the phrase was found in our dataset:

AI Text	**12.44 / 1,000,000 Documents**
Human Text	**1.55 / 1,000,000 Documents**

6x **Reinforcement Learning (via**

How frequently the phrase was found in our dataset:

AI Text	**4.92 / 1,000,000 Documents**
Human Text	**0.77 / 1,000,000 Documents**

6x **matches the historical**

How frequently the phrase was found in our dataset:

AI Text	**2.23 / 1,000,000 Documents**
Human Text	**0.38 / 1,000,000 Documents**

5x **rendering the model**

How frequently the phrase was found in our dataset:

AI Text	**4.42 / 1,000,000 Documents**
Human Text	**0.82 / 1,000,000 Documents**

5x **rigorous mathematical justification**

How frequently the phrase was found in our dataset:

AI Text	**1.91 / 1,000,000 Documents**
Human Text	**0.36 / 1,000,000 Documents**

5x **the RL policy**

How frequently the phrase was found in our dataset:

AI Text	**5.2 / 1,000,000 Documents**
Human Text	**1 / 1,000,000 Documents**

5x **the RL policy**

How frequently the phrase was found in our dataset:

AI Text	**5.2 / 1,000,000 Documents**
Human Text	**1 / 1,000,000 Documents**

5x **as the model has**

How frequently the phrase was found in our dataset:

AI Text	**7.4 / 1,000,000 Documents**
Human Text	**1.46 / 1,000,000 Documents**

5x **bias" stemming from**

How frequently the phrase was found in our dataset:

AI Text	**6.64 / 1,000,000 Documents**
Human Text	**1.44 / 1,000,000 Documents**

4x **in sepsis management.**

How frequently the phrase was found in our dataset:

AI Text	**1.63 / 1,000,000 Documents**
Human Text	**0.37 / 1,000,000 Documents**

4x **Space: The continuous**

How frequently the phrase was found in our dataset:

AI Text	**2.42 / 1,000,000 Documents**
Human Text	**0.57 / 1,000,000 Documents**

4x **(Doubly Robust estimation),**

How frequently the phrase was found in our dataset:

AI Text	**4.1 / 1,000,000 Documents**
Human Text	**0.98 / 1,000,000 Documents**

4x **to detect or correct**

How frequently the phrase was found in our dataset:

AI Text	**1.54 / 1,000,000 Documents**
Human Text	**0.43 / 1,000,000 Documents**

3x **methodological flaws, and**

How frequently the phrase was found in our dataset:

AI Text	**5 / 1,000,000 Documents**
Human Text	**1.46 / 1,000,000 Documents**

3x **Layer to compute**

How frequently the phrase was found in our dataset:

AI Text	**3.72 / 1,000,000 Documents**
Human Text	**1.16 / 1,000,000 Documents**

3x **high performance within**

How frequently the phrase was found in our dataset:

AI Text	**2.24 / 1,000,000 Documents**
Human Text	**0.7 / 1,000,000 Documents**

3x **relies directly on the**

How frequently the phrase was found in our dataset:

AI Text	**1.62 / 1,000,000 Documents**
Human Text	**0.54 / 1,000,000 Documents**

Summary:

The paper proposes SepsisAgent, an agentic Large Language Model (LLM) framework designed for decision support in sepsis management. The method utilizes a three-stage curriculum: Guideline-Aligned Supervision, Tool-Use Behavioral Cloning, and Agentic Reinforcement Learning (via GRPO). A core component of this framework is a GRU-based 'Clinical World Model' that simulates patient physiological responses, allowing the LLM to perform a predict-simulate-refine loop before prescribing interventions. While the integration of LLM reasoning with a world simulator for medical decision-making is conceptually ambitious, the empirical validation suffers from severe methodological flaws, and the underlying Markov Decision Process (MDP) design fundamentally misaligns with clinical reality.

Strengths And Weaknesses:

Strength

Novelty in Pipeline Design: The three-stage training curriculum logically transitions the LLM from basic clinical reasoning to agentic tool use and finally to strategic planning.

Neuro-symbolic Attempt: Incorporating a differentiable Soft Logic Layer to compute clinical scores (SOFA/SIRS) ensures that intermediate simulator outputs are grounded in recognized medical metrics.

Weakness

Evaluation Tautology and Self-Fulfilling Prophecy The primary empirical claim—that SepsisAgent achieves a superior survival rate (59.75%)—is fundamentally flawed due to circular evaluation. The main 'Model-Based Evaluation' is conducted as a roll-out within the exact same Clinical World Model environment used to train the Stage III RL agent. Because the RL policy (GRPO) was explicitly optimized against this specific simulator's dynamics, achieving high performance within it is a tautological result.
Furthermore, on the sole objective, off-policy evaluation metric (Doubly Robust estimation), SepsisAgent (9.37) is outperformed by the traditional AI Clinician baseline (10.47). The authors attribute this to 'systematic bias' stemming from low action matching, but this explanation lacks rigorous mathematical justification and reads as an ad-hoc dismissal of an unfavorable result.

Vulnerability to Reward Hacking and OOD Exploitation Model-based RL approaches are notoriously susceptible to compounding simulator errors. In this work, the World Model is a relatively shallow 2-layer GRU. When the RL policy explores out-of-distribution (OOD) actions, the GRU is highly likely to hallucinate smooth, biologically implausible physiological trajectories. Because the intermediate reward signal relies directly on the predicted changes in SOFA and lactate from this simulator, the agent is effectively optimizing a 'false reward' during hallucinated states. Relying purely on semantic 'clinical skepticism' in the LLM prompt is mathematically insufficient to detect or correct quantitative, dynamic system errors generated by the GRU.

Severe Misalignment Between MDP Formulation and Clinical RealityDespite using a highly capable LLM, the environment formulation severely restricts its potential and contradicts clinical guidelines for sepsis:

(1) Temporal Coarseness: Sepsis is an acute condition where hemodynamics can crash in minutes. Aggregating states into 4-hour timesteps

entirely obscures acute physiological deterioration, rendering the model useless for realistic critical care responses.

(2) Toy Action Space: The continuous spectrum of IV fluids and vasopressors is discretized into a rigid grid. Utilizing an LLM with advanced reasoning capabilities only to output a GridWorld-style discrete action represents a severe under-utilization of the model's capacity for fine-grained dosage adjustments.

Hindsight Bias and 'Teacher Forcing' in Data Generation: The prompt template for Stage I and II data generation (Figure 7) explicitly instructs the teacher LLM to: 'Ensure your reasoning logically leads to the 'Target Action'.' This induces a massive hindsight bias. The SFT data does not represent genuine causal clinical reasoning; rather, it forces the LLM to fabricate a post-hoc rationalization that perfectly matches the historical action. This directly explains the artificially high Action Match rate (31.95%) achieved in Stage II, as the model has merely learned to reverse-engineer justifications for a predetermined target.